

CS289 - Homework 1 Writeup

Suchismita Padhy

Kaggle : Suchi

January 30, 2017

1 Introduction

In this homework we implemented classification using SVM for different datasets such as MNIST digits, CIFAR-10 Images, spam-ham emails.

2 Data Partitioning

- MNIST: We were provided with 2 .mat files for training and test purposes. In order to split the data into training and validation, train data was randomly permuted. First 10,000 rows were held out for validation and the rest was used for training.
- Spam-Ham: We were provided with a directory containing text emails belonging to spam, ham and test. After running the featurize.py code we were able to generate the features and get the mat file. We loaded this mat file which had both train and test data. In order to split the data into training and validation, train data was randomly permuted. 20% of the data was held out for validation and the rest was used for training.
- CIFAR: We were provided with 2 .mat files for training and test purposes. In order to split the data into training and validation, train data was randomly permuted. First 5000 rows were held out for validation and the rest was used for training.

3 Training a Linear SVM

- MNIST:: The SVM classifier was trained for different number of training examples. We used a linear kernel with default values of c .
Below is a plot showing the error rate on the training and validation sets versus the number of training examples that was used to train the classifier.




Figure 1: MNIST: No. of Training Examples Vs Error Rates

- Spam-Ham: The SVM classifier was trained for different number of training examples (100,200,500,1000,2000). We could not use 5000/10,000 as our training data had only 4137 examples. We used a linear kernel with default values of c .

Below is a plot showing the error rate on the training and validation sets versus the number of training examples that was used to train the classifier.




Figure 2: SPAM: No. of Training Examples Vs Error Rates

- CIFAR-10: The SVM classifier was trained for different number of training examples (100,200,500,1000,2000,5000). We used a linear kernel with

default values of c .

Below is a plot showing the error rate on the training and validation sets versus the number of training examples that was used to train the classifier.




Figure 3: CIFAR: No. of Training Examples Vs Error Rates

4 Hyperparameter Tuning

- MNIST: In this step the model was trained with different values of C and the corresponding accuracies were noted. Finally, the C value which gave the highest accuracy on the validation dataset was chosen.

Different Values of C	Accuracies
1e-9	66.26
1e-8	88.49
1e-7	91.55
1e-6	92.34
1e-5	91.15
1e-4	90.4
1e-3	90.4
1e-2	90.4
1e-1	90.4

Table 1: MNIST Hyperparameter Tuning

From the table: 1 we can see that the best accuracy is obtained when C value is between $1e-7$ to $1e-5$. Hence, that range was further explored and the accuracies are shown in table:2.

[0.9193, 0.9208, 0.9222, 0.9122, 0.9193, 0.9142, 0.9133, 0.9122]

Different Values of C	Accuracies
2e-7	91.93
4e-7	92.08
6e-7	92.22
8e-6	91.22
2e-6	91.93
4e-6	91.42
6e-6	91.33
8e-6	91.22

Table 2: MNIST Hyperparameter further exploration

Finally, we chose $6 * 10^{-7}$ as our optimal value of C.

5 K Fold Cross Validation

- Spam-Ham The table 3 shows the different values of C and the corresponding accuracies by averaging the accuracies for 5 folds of cross validation.

Different Values of C	Mean Accuracies
10	80.72
20	80.72
30	80.69
40	80.70
50	80.72
60	80.74
70	80.74
80	80.73
90	80.73
100	80.73

Table 3: SPAM Hyperparameter Tuning using Cross Validation

6 Kaggle

For the kaggle submission, more features were added to increase the score.

- MNIST: I tried adding features using HOG feature engineering method which did not improve my accuracy.
- Spam-Ham: I added a few custom features in featurize.py. Inorder to find these custom features I found frequencies of all the words that were spam and ham and sorted them. I used some of the most frequently occurring

words. I also tried to create my own features using tf-idf. I also tried using different kernels such as polynomial kernels (degree=2,3) and rbf kernel with different values of gamma and C. After these changes the final accuracy rose to 99.13% on cross validation.

However, on the kaggle submission the accuracy was recorded as 93.06%.

7 Appendix : Code

Imports

```
In [1]: import numpy as np
import sklearn as sk
import scipy.io as scio
from sklearn.svm import SVC
from matplotlib import pyplot as plt
import pandas as pd
import os

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import cross_val_score
import re

import skimage as skimg
from skimage.feature import hog
```

Data Loading and Partitioning

```

In [2]: def load_dataset(datasetname):
        location = 'hw01_data/'
        if datasetname == 'mnist':
            location = location + datasetname + '/'
            train_data = scio.loadmat(location + 'train.mat') #load .mat files in python
            test_data = scio.loadmat(location + 'test.mat')

            train_data = train_data['trainX']
            indices = np.random.permutation(len(train_data))

            validation_data = train_data[indices[:10000], :] #First 10000 indices + all columns reserved for validation set
            train_data = train_data[indices[10000:], :] #rest i.e (from 10000th row till end) is train set
            test_data = test_data['testX']

        elif datasetname == 'spam':
            location = location + datasetname + '/'
            spam_data = scio.loadmat(location + 'spam_data.mat')
            train_data = spam_data['training_data']
            test_data = spam_data['test_data']
            train_labels = spam_data['training_labels']

            indices = np.random.permutation(len(train_data))

            split_index = int(0.8 * (len(train_data)))

            validation_data = train_data[indices[split_index:], :]
            train_data = train_data[indices[:split_index], :]

            validation_labels = train_labels[0][indices[split_index:]]
            train_labels = train_labels[0][indices[:split_index]]

            ##Let's just compact train and validation to match with the return type
            train_data = (train_data, train_labels)
            validation_data = (validation_data, validation_labels)

        elif datasetname == 'cifar':
            location = location + datasetname + '/'
            train_data = scio.loadmat(location + 'train.mat')
            test_data = scio.loadmat(location + 'test.mat')

            train_data = train_data['trainX']
            indices = np.random.permutation(len(train_data))

            validation_data = train_data[indices[:5000], :]
            train_data = train_data[indices[5000:], :]
            test_data = test_data['testX']

        return train_data, validation_data, test_data

```

Training a Classifier

```
In [3]: #This function takes the training dataset and C as parameter and returns the trained classifier.

def classify_train(train_data,c=None):
    train_x = train_data[:, :-1]
    train_y = train_data[:, -1]
    if c is not None:
        clf = SVC(kernel='linear', C = c)
    else:
        clf = SVC(kernel="linear")
    clf.fit(train_x, train_y)
    return clf
```

Prediction

```
In [4]: #This function takes the dataset and the trained classifier to return the prediction obtained.

def predict(data,clf):

    data_x = data[:, :-1]
    data_y = data[:, -1]
    predicted_y = clf.predict(data_x)

    return predicted_y
```

MNIST

```
In [5]: mnist_train, mnist_val, mnist_test = load_dataset('mnist')
        # print(mnist_test[0, :])
```

Training on different number of examples


```
In [ ]: iterations = [100,200,500,1000,2000,5000,10000]
val_error_rates = []
train_error_rates = []
classifiers = []
train_y = mnist_train[:, -1]
val_y = mnist_val[:, -1]

for entry in iterations:
    clf = classify_train(mnist_train[:entry, :])

    #train accuracy
    train_predicted_label = predict(mnist_train, clf)
    train_accuracy = len(list(filter(lambda y: y, train_y == train_predicted_label)))/float(len(train_y))
    train_error_rates.append(1-train_accuracy)

    #validation accuracy
    val_predicted_label = predict(mnist_val, clf)
    val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
    val_error_rates.append(1-val_accuracy)
    classifiers.append(clf)
```

Error Plots

```
In [ ]: MNIST_fig = plt.figure()
plt.plot(iterations, list(map(lambda x:x*100,train_error_rates)), 'bo-')
plt.plot(iterations, list(map(lambda x:x*100,val_error_rates)), 'ro-')
plt.legend(["Training Error", "Validation Error"], loc='upper right')
plt.xlabel("No. of Training Examples")
plt.ylabel("Error Rate (in %)")
plt.title("MNIST No. of Training Examples Vs Error Rates")
plt.axis([0, 10500, 0, 100])
plt.show()
MNIST_fig.savefig("MNIST_Training_Examples.png")
```

Finding the best C

```
In [ ]: error_rates = []
classifiers = []
accuracies = []
# c_values = [1e-9,1e-8,1e-7,1e-6,1e-5,1e-4,1e-3,1e-2,1e-1] #best for 1e-6
#Accuracies : [0.6626, 0.8849, 0.9155, 0.9234, 0.9115, 0.904, 0.904, 0.904, 0.904]

# c_values = [2e-7,4e-7,6e-7,8e-6,2e-6,4e-6,6e-6,8e-6] #best for 6e-7
#accuracies = [0.9193, 0.9208, 0.9222, 0.9122, 0.9193, 0.9142, 0.9133, 0.9122]

for c in c_values:
    clf = classify_train(mnist_train[:10000,:],c)

    #validation accuracy

    val_predicted_label = predict(mnist_val, clf)
    val_y = mnist_val[:,-1]
    val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
    error_rates.append(1-val_accuracy)
    accuracies.append(val_accuracy)
    classifiers.append(clf)
print(accuracies)
```

```

In [ ]: # from sklearn.model_selection import GridSearchCV
# from time import time

# # Utility function to report best scores
# def report(results, n_top=3):
#     for i in range(1, n_top + 1):
#         candidates = np.flatnonzero(results['rank_test_score'] == i)
#         for candidate in candidates:
#             print("Model with rank: {0}".format(i))
#             print("Mean validation score: {0:.3f} (std: {1:.3f})".format(
#                 results['mean_test_score'][candidate],
#                 results['std_test_score'][candidate]))
#             print("Parameters: {0}".format(results['params'][candidate]))
#             print("")

# train_x = mnist_train[:2000, :-1]
# train_y = mnist_train[:2000, -1]
# param_grid = [
#     {'C': [1e-1, 1e-2, 1e-3, 1e0], 'kernel': ['linear']},
#     {'C': [1e-1, 1e-2, 1e-3, 1e0], 'gamma': [0.001, 0.0001], 'kernel':
#      ['rbf']},
# ]
# clf = SVC()
# grid_search = GridSearchCV(clf, param_grid=param_grid)
# start = time()
# grid_search.fit(train_x, train_y)

# print("GridSearchCV took %.2f seconds for %d candidate parameter settings."
#       % (time() - start, len(grid_search.cv_results_['params'])))
# report(grid_search.cv_results_)

```

```

In [6]: train_x = mnist_train[:, :-1]
train_y = mnist_train[:, -1]
clf = SVC(C=6e-7, kernel='linear', verbose=True)
clf.fit(train_x, train_y)
val_x = mnist_val[:, :-1]
val_y = mnist_val[:, -1]
val_predicted_label = clf.predict(val_x)
val_accuracy = len(list(filter(lambda y: y, val_y ==
val_predicted_label)))/float(len(val_y))
print(val_accuracy)

```

[LibSVM]0.9461

For the final run let's combine train and validation and train the classifier on the whole data

```

In [7]: train_x = np.vstack([mnist_train[:, :-1], mnist_val[:, :-1]])
train_y = np.concatenate([mnist_train[:, -1], mnist_val[:, -1]])

```

```
In [8]: #Using the parameters that we found to be the best
        clf = SVC(C=6e-7, kernel='linear', verbose=True)
        clf.fit(train_x,train_y)

        [LibSVM]
```

```
Out[8]: SVC(C=6e-07, cache_size=200, class_weight=None, coef0=0.0,
           decision_function_shape=None, degree=3, gamma='auto', kernel='linear',
           max_iter=-1, probability=False, random_state=None, shrinking=True,
           tol=0.001, verbose=True)
```

```
In [9]: # selected_clf = classifiers[np.argmin(error_rates)]

        predicted_label = clf.predict(mnist_test)
```

```
In [10]: output = pd.DataFrame(predicted_label)
         output.columns = ["Category"]
         output.index.names = ["Id"]
         output.to_csv(path_or_buf="submission_mnist.csv", sep=",")
```

Experiments with HOG and SIFT to improve kaggle score

```
In [ ]: train_ims = np.reshape(train_x,[60000,28,28])
        test_ims = np.reshape(mnist_test,[10000,28,28])
```

```
In [ ]: train_features=hog(train_ims[0,:,:])
        for i in range(1,60000):
            train_features = np.vstack([train_features,hog(train_ims[i,:,:])])
            if i%5000==0:
                print(i,end=": ")
                print(train_features.shape)
        train_features.shape
```

```
In [ ]: test_features=hog(test_ims[0,:,:])
        for i in range(1,10000):
            test_features = np.vstack([test_features,hog(test_ims[i,:,:])])
        test_features.shape
```

```
In [ ]: clf = SVC(C=1e-4, kernel='linear')
        scores = cross_val_score(clf, train_features, train_y, cv=2)
        print(scores)
```

```
In [ ]: clf.fit(train_features, train_y)
```

```
In [ ]: predicted_label = clf.predict(test_features)
        output = pd.DataFrame(predicted_label)
        output.columns = ["Category"]
        output.index.names = ["Id"]
        output.to_csv(path_or_buf="submission_mnist.csv", sep=",")
```

CIFAR

```
In [ ]: cifar_train, cifar_val, cifar_test = load_dataset('cifar')
```

Training on different number of examples

```
In [ ]: iterations = [100,200,500,1000,2000,5000]
train_error_rates = []
val_error_rates = []
classifiers = []
train_y = cifar_train[:, -1]
val_y = cifar_val[:, -1]
for entry in iterations:
    clf = classify_train(cifar_train[:entry, :])

    #train accuracy
    train_predicted_label = predict(cifar_train, clf)
    train_accuracy = len(list(filter(lambda y: y, train_y == train_predicted_label)))/float(len(train_y))
    train_error_rates.append(1-train_accuracy)

    #validation accuracy
    val_predicted_label = predict(cifar_val, clf)
    val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
    val_error_rates.append(1-val_accuracy)
    print("done : " + str(entry))
    classifiers.append(clf)
```

Error plots

```
In [ ]: CIFAR_fig = plt.figure()
plt.plot(iterations, list(map(lambda x:x*100,train_error_rates)), 'bo-')
plt.plot(iterations, list(map(lambda x:x*100,val_error_rates)), 'ro-')
plt.legend(["Training Error", "Validation Error"], loc='upper right')
plt.xlabel("No. of Training Examples")
plt.ylabel("Error Rate (in %)")
plt.title("CIFAR No. of Training Examples Vs Error Rates")
plt.axis([0, 5500, 0, 100])
plt.show()
CIFAR_fig.savefig("CIFAR_Training_Examples.png")
```

Finding the best C

```
In [ ]: error_rates = []
classifiers = []
c_values = [1e-4,1e-3,1e-2,1e-1,1e0,1e1,1e2,1e3]
for c in c_values:
    clf = classify_train(cifar_train[:1000,:],c)

    #validation accuracy
    val_predicted_label = predict(cifar_val, clf)
    val_y = cifar_val[:,-1]
    val_accuracy = len(list(filter(lambda y: y, val_y != val_predicted_label))) / float(len(val_y))
    error_rates.append(1-val_accuracy)
    classifiers.append(clf)
print(error_rates)
```

```
In [ ]: train_x = cifar_train[:,-1]
train_y = cifar_train[:,-1]
clf = SVC(kernel='poly', C =1e-3, degree=2)
clf.fit(train_x,train_y)
val_x = cifar_val[:,-1]
val_y = cifar_val[:,-1]
val_predicted_label = clf.predict(val_x)
val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label))) / float(len(val_y))
print(val_accuracy)
```

```
In [ ]: train_x = cifar_train[:,-1]
train_y = cifar_train[:,-1]
print(train_y.shape)
```

```
In [ ]: # selected_clf = classifiers[np.argmin(error_rates)]
predicted_label = clf.predict(cifar_test)
predicted_label
```

Spam - Ham

```
In [11]: spam_train, spam_val, spam_test = load_dataset('spam')
```

```
In [12]: train_x,train_y = spam_train
val_x,val_y= spam_val
len(train_x)
```

```
Out[12]: 4137
```

Training on different number of examples

```
In [ ]: iterations = [100,200,500,1000,2000]

val_error_rates = []
train_error_rates = []

classifiers = []

for entry in iterations:
    clf = SVC(kernel="linear")
    clf.fit(train_x[:entry,:],train_y[:entry])
    #train accuracy
    train_predicted_label = clf.predict(train_x)
    train_accuracy = len(list(filter(lambda y: y, train_y == train_predicted_label)))/float(len(train_y))
    train_error_rates.append(1-train_accuracy)

    #validation accuracy
    val_predicted_label = clf.predict(val_x)
    val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
    val_error_rates.append(1-val_accuracy)
    classifiers.append(clf)
```

Error plot

```
In [ ]: Spam_fig = plt.figure()
plt.plot(iterations, list(map(lambda x:x*100,train_error_rates)),'bo-')
plt.plot(iterations, list(map(lambda x:x*100,val_error_rates)),'ro-')
plt.legend(["Training Error","Validation Error"], loc='upper right')
plt.xlabel("No. of Training Examples")
plt.ylabel("Error Rate (in %)")
plt.title("SPAM-HAM No. of Training Examples Vs Error Rates")
plt.axis([0, 2500,0,100])
plt.show()
Spam_fig.savefig("SPAM_Training_Examples.png")
```

```
In [ ]: # clf = SVC(kernel='rbf',gamma=0.8e-2,C=55)
# clf.fit(train_x,train_y)
# val_predicted_label = clf.predict(val_x)
# val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
# print(val_accuracy)
```

```
In [ ]: clf = SVC(kernel='linear',C=55)
clf.fit(train_x,train_y)
val_predicted_label = clf.predict(val_x)
val_accuracy = len(list(filter(lambda y: y, val_y == val_predicted_label)))/float(len(val_y))
print(val_accuracy)
```

K fold Cross Validation

```
In [ ]: train_x = np.vstack([train_x, val_x])
train_y = np.concatenate([train_y, val_y])
C_values = list(range(10, 200, 10))
scores = []
for c in C_values:
    clf = SVC(kernel='linear', C=c)
    score = cross_val_score(clf, train_x, train_y, cv=5)
    avg_score = np.mean(score)
    scores.append(avg_score)
    print(c, end=": ")
    print(avg_score)
```

Trying TfIdf for better score

```
In [13]: spam_dir = "hw01_data/spam/spam/"
ham_dir = "hw01_data/spam/ham/"
test_dir = "hw01_data/spam/test/"
spam_files = filter(lambda x: x.endswith('.txt'), os.listdir(spam_dir))
ham_files = filter(lambda x: x.endswith('.txt'), os.listdir(ham_dir))
test_files_len = len(list(filter(lambda
x: x.endswith('.txt'), os.listdir(test_dir))))
```

```
In [14]: df = pd.DataFrame(columns=("text", "label")) # create an empty dataframe
count=0
for fn in spam_files:
    fn=spam_dir+fn
    with open(fn, "r", encoding='utf-8', errors='ignore') as spam:
        df.loc[count] = [spam.read(), 1]
        count+=1

for fn in ham_files:
    fn=ham_dir+fn
    with open(fn, "r", encoding='utf-8', errors='ignore') as ham:
        df.loc[count] = [ham.read(), 0]
        count+=1
df = df.sample(frac=1).reset_index(drop=True) # shuffle the dataframe
```

```
In [15]: # df_val = df.ix[:1000]
# df_train = df.ix[1001:]
```



```
In [16]: #do the same thing for test dataset
df_test = pd.DataFrame(columns=("text","label"))
count=0
for fn in range(test_files_len):
    fn=test_dir+str(fn)+'.txt'
    with open(fn,"r",encoding='utf-8', errors='ignore') as test:
        df_test.loc[count] = [test.read(),"UNK"]
        count+=1
```

```
In [17]: #Tfidf
df_train = df
vectorizer = TfidfVectorizer(sublinear_tf=True, stop_words='english')
X_train = vectorizer.fit_transform(df_train.text)
# X_val = vectorizer.transform(df_val.text)
X_test = vectorizer.transform(df_test.text)
```

```
In [18]: #classifier

# clf = SVC(kernel='linear')
clf = SVC(kernel='rbf',gamma=0.002,C=300)
scores = cross_val_score(clf, X_train, df_train.label, cv=5)
print(scores)
clf.fit(X_train,df_train.label)
# val_predicted_label = clf.predict(X_val)
# val_accuracy = len(list(filter(lambda y: y, df_val.label == val_predicted_label)))/float(len(df_val))
# print(val_accuracy)

[ 0.99130435  0.99130435  0.98839458  0.98839458  0.98646035]
```

```
Out[18]: SVC(C=300, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape=None, degree=3, gamma=0.002, kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

```
In [19]: predicted_label = clf.predict(X_test)
output = pd.DataFrame(predicted_label)
output.columns = ["Category"]
output.index.names = ["Id"]
output.to_csv(path_or_buf="submission_spam.csv", sep=",")
```

Checking the frequency of words to decide optimum features to be used in featurizer.py

```
In [ ]: # def tokenizer(text):
#         return re.split(" |\n",text)
# df_spam = df_train[df_train.label==0]
```

```
In [ ]: # df_spam=df_spam.reset_index()
```

```
In [ ]: # corpus ={}  
# for i in range(len(df_spam)):  
#     for token in tokenizer(df_spam.text[i]):  
#         corpus[token] = corpus.get(token,0)+1
```

```
In [ ]: # import operator  
# sorted(corpus.items(), key=operator.itemgetter(1),reverse=True)
```

```
In [ ]:
```